



2025 Extended Report

<https://chc-comp.github.io/>

presented at
HCVS 2025, July 22, Zagreb, Croatia

Gidon Ernst, LMU Munich, Germany

Jose F. Morales, IMDEA Software Institute, Spain

Levente Bajczi, Budapest University of Technology and Economics

Goals & Overview

- CHC-COMP: friendly but competitive evaluation of constrained Horn-clause solvers, since 2018
<https://chc-comp.github.io/>
- common task format (subset of SMT-LIB)
<https://chc-comp.github.io/format.html>
- public benchmark repository (please submit!)
<https://github.com/chc-comp>
- Timeline: Jan—May, results presented (at SPIN and) HCVS

Setup and Updates in 2025



- Move from StarExec to LMU cluster (SV-COMP infrastructure)
 - Intel Xeon E3-1230 v5 @ 3.40 GHz, 8 cores, 30 GB memory, 1800s
 - <https://github.com/chc-comp/chc-comp25-scripts>
- Add **BV**, **LRA-Lin** tracks, evaluate **all feasible benchmarks**
 - <https://github.com/chc-comp/chc-comp25-benchmarks>
 - overall CPU time: ~half a year (including some trial&error)
- Model validation done for HCVS@CAV 2025 (**thanks to Levente**)
- **thanks to Dirk Beyer and members of SoSy Lab for support**

2025: Participants

1. CHC2C 1.0 (Mihály Dobos-Kovács, Levente Bajczi, András Vörös),
Note: meta-solver
2. ChocoCatalia (Hiroyuki Katsura, Naoki Kobayashi, Ryosuke Sato)
3. [Eldarica 2.2](#) (Hossein Hojjat, Philipp Ruemmer)
4. [Golem 0.7.1](#) (Martin Blicha)
5. [LoAT](#) (Florian Frohn, Jürgen Giesl)
6. [MuCyc](#) (Kazuki Uehara, Hiroshi Unno)
7. [PCSat](#) (Takuma Monma, Hiroshi Unno)
8. [ThetaCHC 6.13.2](#) (Levente Bajczi, Mihály Dobos-Kovács, Márk Somorjai, András Vörös)
9. [Ultimate Tree Automizer](#) (Matthias Heizmann, Max Barth, Daniel Dietsch, Dominik Klumpp)
10. [Ultimate Unihorn](#) (Matthias Heizmann, Max Barth, Daniel Dietsch, Dominik Klumpp)

	LIA- Lin 1312	LIA 1266	LIA- Lin- Arrays 139	LIA- Arrays 1728	ADT- LIA 3585	ADT- LIA- Arrays 1045	BV 559	LRA- Lin 274
<i>CHC2C (meta)</i>	✓	✓						
ChocoCatalia					✓			
MuCyC	✓	(*)						
PCSat	✓	✓	✓	✓	✓	✓	✓	✓
Eldarica	✓	✓	✓	✓	✓	✓	✓	✓
Golem	✓	✓	✓					✓
LoAT	✓							
Theta	✓	✓	✓	✓			✓	✓
U. Tree Automizer	✓	✓	✓	✓				
Ultimate Unihorn	✓	✓	✓	✓				

Z3/Spacer was not submitted

(*) I somehow missed running this (apologies!)

2025: Winners*

LIA-Lin	LIA	LIA-Lin- Arrays**	LIA- Arrays	ADT- LIA	ADT-LIA- Arrays	BV	LRA-Lin
---------	-----	----------------------	----------------	-------------	--------------------	----	---------

Golem	Golem	Eldarica	Eldarica	Catalia	Eldarica	Eldarica	Golem
-------	-------	----------	----------	---------	----------	----------	-------

MuCyc	Eldarica	Unihorn	PCSat	Eldarica	PCSat	Theta	Eldarica
-------	----------	---------	-------	----------	-------	-------	----------

LoAT	PCSat	PCSat	Unihorn	PCSat	--	PCSat	Theta
------	-------	-------	---------	-------	----	-------	-------



*accepting all results as correct
even though there are definitely inconsistencies

**ranking changed after cleanup of results

2025: Infrastructure

<https://gitlab.com/sosy-lab/software/benchcloud>



benchexec: resource control, measurements [Beyer+ STTT19]

- Python toolinfo modules as API for tools (provided by organizer)
Examples: **chc.py** and **eldarica.py**
- XML benchmark definition files to set up experiments (provided)
Example: **golem.xml**

Should establish community process to maintain these s part of **FM tools** |
(→ akin to SV-COMP but keep it light-weight!)

benchcloud: cluster management, job scheduling [Beyer+ ASE24]

- deployed at LMU and also on a Hungarian cluster (this helped!)

Pitfall: duplicate benchmark names not supported (→ design decision)

2025: Scripts

<https://github.com/chc-comp/chc-comp25-scripts>

Tool-chain accumulated over past years

- **format.py** ensure CHC-COMP requirements, obfuscation
(has several problems, should be replaced next year)
- **classify.py** determine benchmark's categories
(new! replaces previous slow/inflexible format-checker)
- scripts to set up benchexec metadata .yaml files
→ cross-check result, found and fixed inconsistencies (bugs)
- **Lacking: fully documented pipeline**

2025: Benchmarks

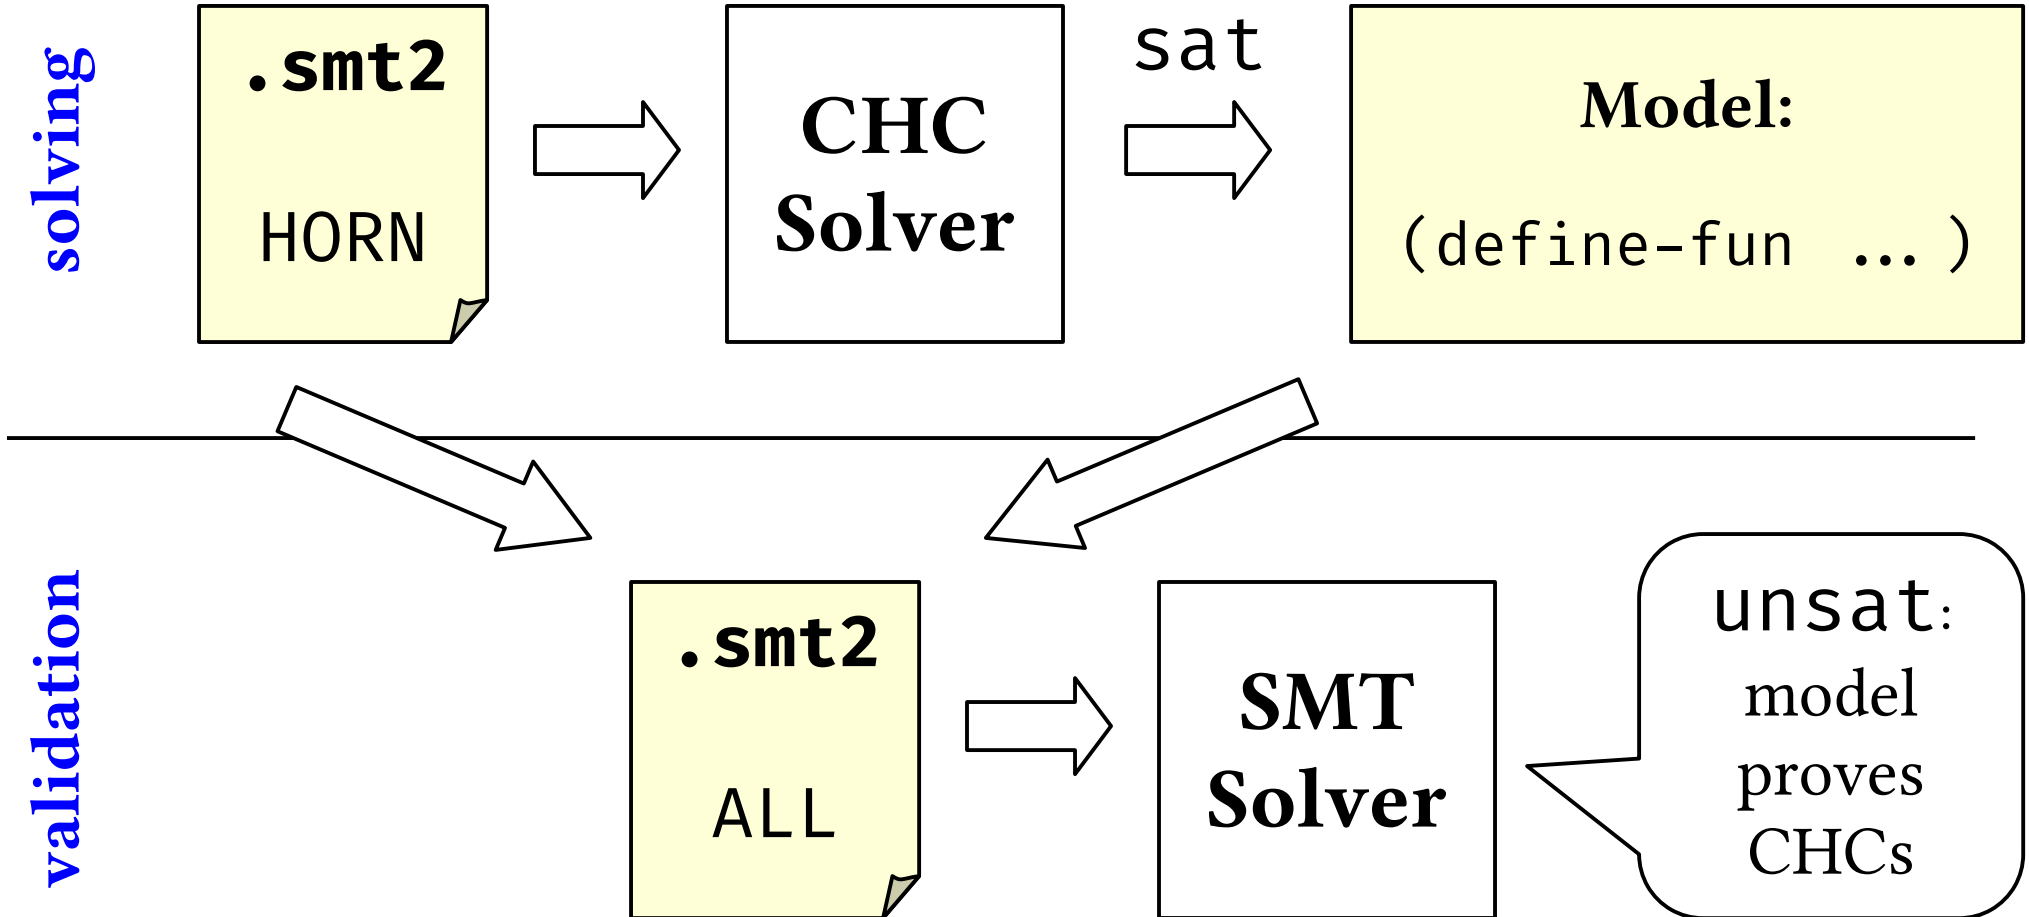
<https://github.com/chc-comp/chc-comp25-benchmarks>

Benchmarks accumulated over past years

- fairly diverse in features and in difficulty
- organization into individual sub-repositories with “raw” sources
- not “ready to use”: format inconsistencies, lack of ground truth
- unclear which benchmarks are “interesting”
- preprocessing pipeline fails on 2/3 of files (!?)

Please contribute benchmarks!

2025: Validation (`validate-model.py`)



2025: Model Validation Results

<https://leventebajcsi.com/chc-model-artifact/>

Category	<u>Eldarica</u>	Golem	<u>ThetaCHC</u>
LIA	378	709	0
LIA-Lin	623	675	565
LIA-Arrays	1000	-	0
LIA-Lin-Arrays	52	0	55
LRA-Lin	0	73	11
BV	17	-	22

points are awarded for correct results (wrt. ground-truth)
for which at least one SMT solver confirms the model

2025: Model confirmation (36%-100%)

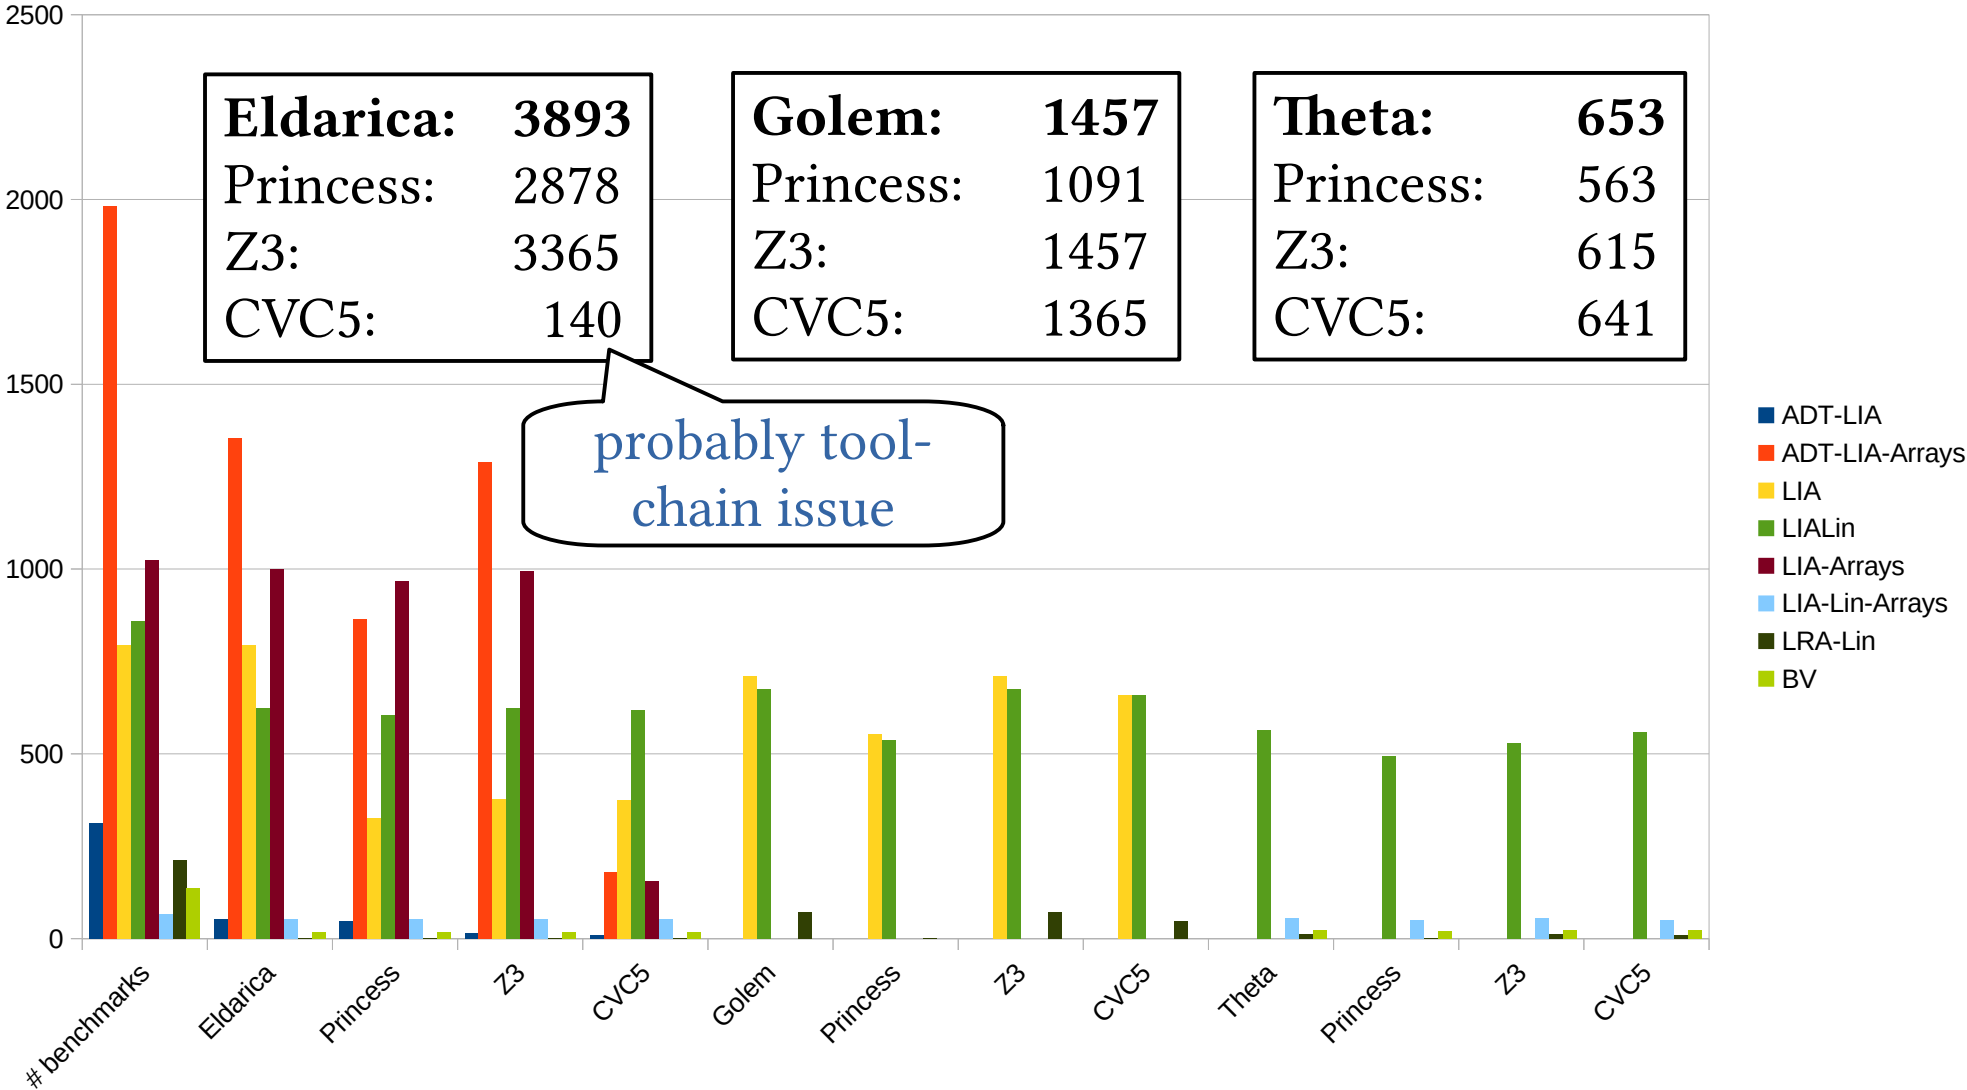
<https://leventebaiczi.com/chc-model-artifact/>

Eldarica:	3893
Princess:	2878
Z3:	3365
CVC5:	140

Golem:	1457
Princess:	1091
Z3:	1457
CVC5:	1365

Theta:	653
Princess:	563
Z3:	615
CVC5:	641

probably tool-chain issue



Summary and Outlook



2025 Achievements (modulo small hick-ups and lessons learned)

- infrastructure switch StarExec → benchcloud
- new benchmarks (LRA, BV)
- first steps towards model validation

Follow-up

- document scripts and pipeline, write a report (with fixed results?)
- discuss benchmark repository (ground truth, pre-processing, file names)
- counterexample proof format and validation
- massively parallel track? “best-effort” track without goal clauses?

Partitipants: publish your archive on zenodo, make an entry in [FM-Tools](#)

Organizers of the next edition? I’m happy to help with onboarding :)